

STAT 542: Statistical Learning

K -Nearest Neighbor and the Bias-Variance Trade-Off

Ruoqing Zhu, Ph.D. <rqzhu@illinois.edu>

Course Website: <https://teazrq.github.io/stat542/>

August 26, 2019

Department of Statistics
University of Illinois at Urbana-Champaign

K-Nearest Neighbour

An accurate prediction

- Let's consider a regression model,

$$Y = f(X) + \epsilon,$$

where $E(\epsilon) = 0$ and $\text{Var}(\epsilon) = \sigma^2$.

- Suppose that we collect a set of training data $\mathcal{D}_n = \{x_i, y_i\}_{i=1}^n$
- from the training data, we are able to estimate the regression function as \hat{f} (“ f -hat”).
- Want to predict the value of Y at a target point x .

An accurate prediction

- Using squared-error loss, the error of the prediction is

$$\begin{aligned}\text{Err}(x) &= \mathbb{E}_{\mathcal{D}, Y} \left[(Y - \hat{f}(x))^2 \right] \\ &= \mathbb{E}_{\mathcal{D}, Y} \left[(Y - f(x) + f(x) - \mathbb{E}_{\mathcal{D}} \hat{f}(x) + \mathbb{E}_{\mathcal{D}} \hat{f}(x) - \hat{f}(x))^2 \right] \\ &= \dots \\ &= \underbrace{\mathbb{E}_Y \left[(Y - f(x))^2 \right]}_{\text{Irreducible Error}} + \underbrace{\left(f(x) - \mathbb{E}_{\mathcal{D}} \hat{f}(x) \right)^2}_{\text{Bias}^2} + \underbrace{\mathbb{E}_{\mathcal{D}} \left[(\hat{f}(x) - \mathbb{E}_{\mathcal{D}} \hat{f}(x))^2 \right]}_{\text{Variance}}\end{aligned}$$

- All the cross terms are zero

An accurate prediction

- $E[(Y - f(x))^2] = \sigma^2$ is the **irreducible error** term that cannot be avoided, because we cannot predict ϵ
- $E[(f(x) - E\hat{f}(x))^2]$ is the **squared bias** term that evaluates how the average of our estimator deviates from the truth
- $E[(\hat{f}(x) - E\hat{f}(x))^2]$ is the **variance** term that reflects the sensitivity of the function estimate $\hat{f}(x)$ to the training sample

Of course we want to minimize both the **Bias²** and the **Variance**, however, this is not always possible...

An accurate prediction

- The first instinct is that one should minimize the bias² term $E[(f(x) - E\hat{f}(x))^2]$
- Why would I produce a model that is asymptotically incorrect, i.e., $f(x) \neq E\hat{f}(x)$?
- **This instinct is wrong!** — This is usually at the expense of high variance, which eventually damages the prediction performance...
- To demonstrate this, let's start with a particular model, the k -Nearest Neighbour, in both classification and regression settings

k -Nearest Neighbour

- k -Nearest Neighbour (k NN) is a nonparametric method that predicts the target point x with averages of nearby observations in the training data
- For regression, the prediction at a given target point x is

$$\hat{y} = \frac{1}{k} \sum_{x_i \in N_k(x)} y_i,$$

where $N_k(x)$ defines the k samples from the training data (in terms of their feature values) that are closest to x .

- For hard classification, the most prevalent class in $N_k(x)$ is used


$$\hat{y} = \arg \max_{c \in C} \sum_{x_i \in N_k(x)} \mathbf{1}\{y_i = c\},$$

k -Nearest Neighbour in Regression

The the following data is observed, with only 1 feature, uniformly from $[0, 2\pi]$. The true model (blue line) is

$$Y = 2 \sin(X) + \epsilon,$$

where ϵ is standard normal error. We fit the data with 1NN.



k -Nearest Neighbour in Regression

- The bias² term is minimized using 1NN, why?

$$\begin{aligned}\mathbb{E}\hat{f}(x) &= \mathbb{E}\left[\sum Y_i \mathbf{1}\{X_i \in N_1(x)\}\right] \\ &= \mathbb{E}\left[\sum f(X_i) \mathbf{1}\{X_i \in N_1(x)\}\right]\end{aligned}$$

which the closest we can get for approximating $f(x)$ (if we believe there is a certain continuity in f).



- However, the variance is large because the estimator only uses one observation

$$\mathbb{E}\left[(\hat{f}(x) - \mathbb{E}\hat{f}(x))^2\right] = \mathbb{E}\epsilon^2 = \sigma^2.$$

- If we use more “neighbouring” points, say k , the variance would reduce to σ^2/k . But the bias² will increase because as the neighbourhood expands, $f(X_i)$ is further away from $f(x)$.



k -Nearest Neighbour in Regression

Now we simulate 200 observations, and see how the model changes over k .



k -Nearest Neighbour in Regression


Now we simulate 200 observations, and see how the model changes over k .




k -Nearest Neighbour in Regression

The model becomes “smoother” as k increases. However, this eventually introduces a significant bias.

33-Nearest Neighbor Regression




100-Nearest Neighbor Regression



Model Complexity, over- and under-fitting

- Model complexity \uparrow (small k) \longrightarrow Bias² \downarrow and Variance \uparrow
- Model complexity \downarrow (large k) \longrightarrow Bias² \uparrow and Variance \downarrow



Prevent over-fitting


- As we can see, model complexity, bias-variance trade-off and over- and under-fitting are usually related concepts
- Over-fitting happens when the model performs well on the training sample, but not on the testing sample
- Controlling complexity can prevent overfitting. Complexity can be measured in different ways. In statistics, we often use degrees of freedom (number of parameters in a model)
- The degrees of freedom of a k NN model is roughly n/k
 - intuition: if neighborhoods don't overlap, there would be n/k neighborhoods, with one parameter for each

Prevent over-fitting

- To control complexity, one approach is to “**penalize**” it (we will see this a lot later on):

$$\arg \min_f \text{loss}(f) + \lambda \text{complexity}(f)$$

- Another common approach is to use **cross-validation** (CV). A 10-fold CV is carried out as follows
 - Randomly split the data into 10 equal sized subsamples
 - Fit the model using 9 out of 10 subsamples as training data and calculate the testing error using the remaining one.
 - Alternate the testing sample, and average the total of 10 experiments



Remark on the degrees of freedom

A rigorous definition for the degrees of freedom is

$$\text{df}(\hat{f}) = \frac{1}{\sigma^2} \sum_{i=1}^n \text{Cov}(\hat{Y}_i, Y_i)$$

- The amount of covariance between outcome Y_i and its fitted values \hat{Y}_i , at the same target point x_i .
- **Treat $X_i = x_i$ as fixed value**, but not random.
- $1/\sigma^2$ takes care of the variance of the random error term.
- In many situations, it is more convenient to define it in the matrix form, noticing that we are just adding up the diagonal elements of the covariance matrix of $\hat{\mathbf{Y}} = (\hat{Y}_1, \dots, \hat{Y}_n)$ and $\mathbf{Y} = (Y_1, \dots, Y_n)$

$$\text{df}(\hat{f}) = \frac{1}{\sigma^2} \text{Trace}(\text{Cov}(\hat{\mathbf{Y}}, \mathbf{Y}))$$

Remark on the degrees of freedom

Based on this definition, we can easily verify several cases:

- For 1NN, $\text{df} = n$
- If $\hat{y}_i = \bar{y}$, i.e., $n\text{NN}$, then $\text{df} = 1$
- For linear regression, $\text{df} = p$
- For $k\text{NN}$, $\text{df} = n/k$

The formula works for kernel methods too, we shall see that later on.

Bias-Variance Trade-off in Classification

- Classification problems have the same bias-variance trade-off, although the formulation is slightly different due to their loss functions (0-1 loss for “hard” classification)
- A very nice paper about this issue by Friedman (1997)
 - “On bias, variance, 0/1—loss, and the curse-of-dimensionality”Try example 7.2 in HTF.
- The optimal classifier is call the Bayes classifier:

$$f_B(x) = \mathbf{1}\{\mathbf{P}(Y|X = x) > 1/2\}$$


- Error produced by Bayes classifier

$$\min \{ \mathbf{P}(Y|X = x), 1 - \mathbf{P}(Y|X = x) \}$$

is analog to the irreducible error σ^2 in a regression setting


k -Nearest Neighbour in Classification

Let's look at a classification example from the HTF text book.
(BLUE = 0, ORANGE = 1)



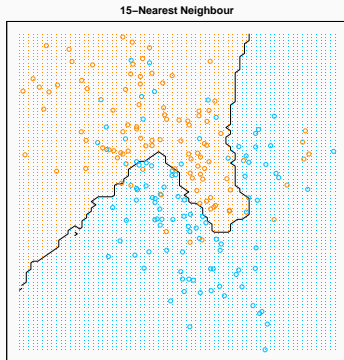
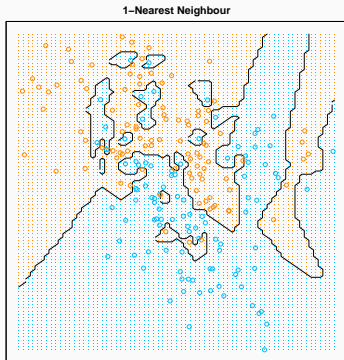
k -Nearest Neighbour in Classification

Similar to the regression case, the k -NN classification model does majority vote (the most prevalent class) within the neighborhood of a target point x . 1NN plot is a Voronoi tessellation





k -Nearest Neighbour in Classification

We fit k -NN classification model to the example. Of course, we would not expect 1NN to perform well...



k -Nearest Neighbour in Classification

As we further increase k , the model tends to be less complex.
Compare 66NN with a linear model that uses only 3 parameters.




k -Nearest Neighbour vs. Linear Regression

- The goal is to approximate $f(x) = E(Y|X = x)$
- Linear regression makes a structural assumption: f is linear.
 - **low variance**: Number of parameters is p (fixed); we know that when sample size n grows, the variance of $\hat{\beta}$ is $\propto 1/n$.
 - **high bias** (underfit): linear assumption is very restrictive
- k NN makes on assumption on f , except some smoothness.
 - **low bias** (overfit): flexible and adaptive. It can be shown that as if $k \rightarrow \infty$ and $n/k \rightarrow 0$, k NN is consistent.
 - **high variance**: number of parameters for k NN is roughly n/k ;

k -Nearest Neighbour in Classification

An “U” shaped prediction error curve is again observed for the testing sample (Figure from HTF):



Distance measures

- Closeness between two points needs to be defined based on some distance measures
- By default, we use Euclidean distance (ℓ_2 norm) for continuous variables

$$d^2(\mathbf{u}, \mathbf{v}) = \|\mathbf{u} - \mathbf{v}\|_2^2 = \sum_{i=1}^p (u_i - v_i)^2$$

Hence the neighbourhood is not invariant to the scaling of the variables.

Scaling issues

- We often scale the variables marginally when using k NN, so that the distance is

$$d^2(\mathbf{u}, \mathbf{v}) = \sum_{j=1}^p \frac{(u_j - v_j)^2}{\sigma_j^2}$$

where σ_j^2 is the variance of variable j .

- Mahalanobis distance is also scale-invariant and takes care of correlation

$$d^2(\mathbf{u}, \mathbf{v}) = (\mathbf{u} - \mathbf{v})^T \Sigma^{-1} (\mathbf{u} - \mathbf{v}),$$

where Σ is a covariance matrix.


- Hamming distance is usually used for categorical variables

Drawbacks of k NN

- Need to score the entire training data for future prediction
- Calculating the prediction can be **slow**. To find the nearest neighbor of x , one needs to calculate the distance from x to all training sample and compare them. This is computationally expansive especially in high-dimensional setting
- There are some fast nearest neighbor search algorithms such as kd-tree
- A distance measure needs to be specified — it may affect accuracy

Example: Handwritten Digit Recognition Data

- Digits 1-9 scanned from envelopes by the U.S. Postal Service
- 7291 training samples, 2007 testing samples
- Apply k NN and calculate the errors
- 1NN with Euclidean distance gives 5.6% error rate
- 1NN with **tangent distance** (Simard et al., 1993) gives 2.6% error




New Challenges

New Challenge

- High-dimension low sample size ($p \gg n$)
 - The resolution of the handwritten digit example is $16 \times 16 = 256$
 - Some common imaging data in medical are 1024×1024 while only a few hundred samples are available
 - Strategy games (Go, StarCraft, etc.) may have a huge number of variables
- Curse of Dimensionality
 - For fixed n , as p increases, the data become sparse
 - As p increases, the number of possible models explodes (computation burden, variable selection necessary)

Curse of Dimensionality


- The curse of dimensionality is well illustrated by a subcubical neighborhood for uniform data in a unit cube.
- Suppose the sample points are evenly spread out on $[0, 1]^p$, and we want to capture 10% of the data by constructing a hypercube neighborhood of x . What is the edge length l of this cube? Since the volume of the cube is $l^p = 10\%$, we need $l = 0.1^{1/p}$,



- When $p = 1$, $l = 0.1$
- When $p = 2$, $l = 0.32$
- When $p = 10$, $l = 0.79$

Curse of Dimensionality

- Suppose we have sample points evenly spread out on $[0, 1]$
- In ten dimensions we need to cover 80% of the range of each coordinate to capture 10% of the data.



Curse of Dimensionality

- In linear regression, the curse of dimensionality is easy to see from the **normal equation**:

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

where $\mathbf{X}_{n \times p}$ is the design matrix, $\mathbf{y}_{n \times 1}$ is the response vector.

- $\mathbf{X}^T \mathbf{X}$ is not invertible when $p > n$, hence $\hat{\beta}$ is not unique
- In fact, since \mathbf{X} is of rank n , we can always find some $\hat{\beta}$ such that $\mathbf{y} = \mathbf{X}\hat{\beta}$, this means we are modeling the error term ϵ !
- Some penalized methods can deal with this problem